



Experiment No. 5
Apply appropriate Unsupervised Learning Technique on the Wholesale Customers Dataset
Date of Performance: 21-08-2023
Date of Submission: 30-09-2023



Aim: Apply appropriate Unsupervised Learning Technique on the Wholesale Customers Dataset.


Objective: Able to perform various feature engineering tasks, apply Clustering Algorithm on the given dataset.

Theory:

It is basically a type of unsupervised learning method. An unsupervised learning method is a method in which we draw references from datasets consisting of input data without labeled responses. Generally, it is used as a process to find meaningful structure, explanatory underlying processes, generative features, and groupings inherent in a set of examples.

Clustering is the task of dividing the population or data points into a number of groups such that data points in the same groups are more similar to other data points in the same group and dissimilar to the data points in other groups. It is basically a collection of objects on the basis of similarity and dissimilarity between them.

For example: The data points in the graph below clustered together can be classified into one single group. We can distinguish the clusters, and we can identify that there are 3 clusters in the below picture.





Dataset:

This data set refers to clients of a wholesale distributor. It includes the annual spending in monetary units (m.u.) on diverse product categories. The wholesale distributor operating in different regions of Portugal has information on annual spending of several items in their stores across different regions and channels. The dataset consist of 440 large retailers annual spending on 6 different varieties of product in 3 different regions (lisbon , oporto, other) and across different sales channel (Hotel, channel)

Detailed overview of dataset

Records in the dataset = 440 ROWS

Columns in the dataset = 8 COLUMNS

FRESH: annual spending (m.u.) on fresh products (Continuous)

MILK:- annual spending (m.u.) on milk products (Continuous)

GROCERY:- annual spending (m.u.) on grocery products (Continuous)

FROZEN:- annual spending (m.u.) on frozen products (Continuous)

DETERGENTS_PAPER :- annual spending (m.u.) on detergents and paper products (Continuous)

DELICATESSEN:- annual spending (m.u.)on and delicatessen products (Continuous);

CHANNEL: - sales channel Hotel and Retailer

REGION:- three regions (Lisbon, Oporto, Other)

```
In [21]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv("customers.csv")
df.sample(5)
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicass
en								
0	2	3	12669	9656	7561	214	2674	13
38								
1	2	3	7057	9810	9568	1762	3293	17
76								
2	2	3	6353	8808	7684	2405	3516	78
44								
3	1	3	13265	1196	4221	6404	507	17
88								
4	2	3	22615	5410	7198	3915	1777	51
85								

```
In [7]: print(df.columns)
```

```
Index(['Channel', 'Region', 'Fresh', 'Milk', 'Grocery', 'Frozen',
      'Detergents_Paper', 'Delicassen'],
      dtype='object')
```

```
In [6]: print(df.dtypes)
```

```
Channel          int64
Region           int64
Fresh            int64
Milk             int64
Grocery          int64
Frozen           int64
Detergents_Paper int64
Delicassen       int64
dtype: object
```

```
In [5]: print(df.isnull().sum())
```

```
Channel          0
Region           0
Fresh            0
Milk             0
Grocery          0
Frozen           0
Detergents_Paper 0
Delicassen       0
dtype: int64
```

```
In [9]: import matplotlib.pyplot as plt
import seaborn as sns

print("Descriptive Statistics:")
print(df.describe())

print("Number of duplicate rows: ", df.duplicated().sum())

for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[column], bins=30, kde=True)
    plt.title(f'Distribution of {column}')
    plt.show()

plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Heatmap')
plt.show()
```


Descriptive Statistics:

	Channel	Region	Fresh	Milk	Grocery \
count	440.000000	440.000000	440.000000	440.000000	440.000000
mean	1.322727	2.543182	12000.297727	5796.265909	7951.277273
std	0.468052	0.774272	12647.328865	7380.377175	9503.162829
min	1.000000	1.000000	3.000000	55.000000	3.000000
25%	1.000000	2.000000	3127.750000	1533.000000	2153.000000
50%	1.000000	3.000000	8504.000000	3627.000000	4755.500000
75%	2.000000	3.000000	16933.750000	7190.250000	10655.750000
max	2.000000	3.000000	112151.000000	73498.000000	92780.000000


	Frozen	Detergents_Paper	Delicassen
count	440.000000	440.000000	440.000000
mean	3071.931818	2881.493182	1524.870455
std	4854.673333	4767.854448	2820.105937
min	25.000000	3.000000	3.000000
25%	742.250000	256.750000	408.250000
50%	1526.000000	816.500000	965.500000
75%	3554.250000	3922.000000	1820.250000
max	60869.000000	40827.000000	47943.000000

Number of duplicate rows: 0


Distribution of Channel




Distribution of Region




Distribution of Fresh




Distribution of Milk




Distribution of Grocery




Distribution of Frozen




Distribution of Detergents_Paper



Distribution of Delicassen





```


In [10]: import seaborn as sns
import matplotlib.pyplot as plt

for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.boxplot(df[column])
    plt.title(f'Boxplot of {column}')
    plt.show()


def detect_outliers(dataframe, column):
    Q1 = dataframe[column].quantile(0.25)
    Q3 = dataframe[column].quantile(0.75)
    IQR = Q3 - Q1
    outliers = dataframe[(dataframe[column] < Q1 - 1.5*IQR) | (dataframe[column] > Q3 + 1.5*IQR)]
    return outliers

for column in df.columns:
    outliers = detect_outliers(df, column)
    print(f'Number of outliers in {column}: {len(outliers)}')


```




Boxplot of Region




Boxplot of Fresh





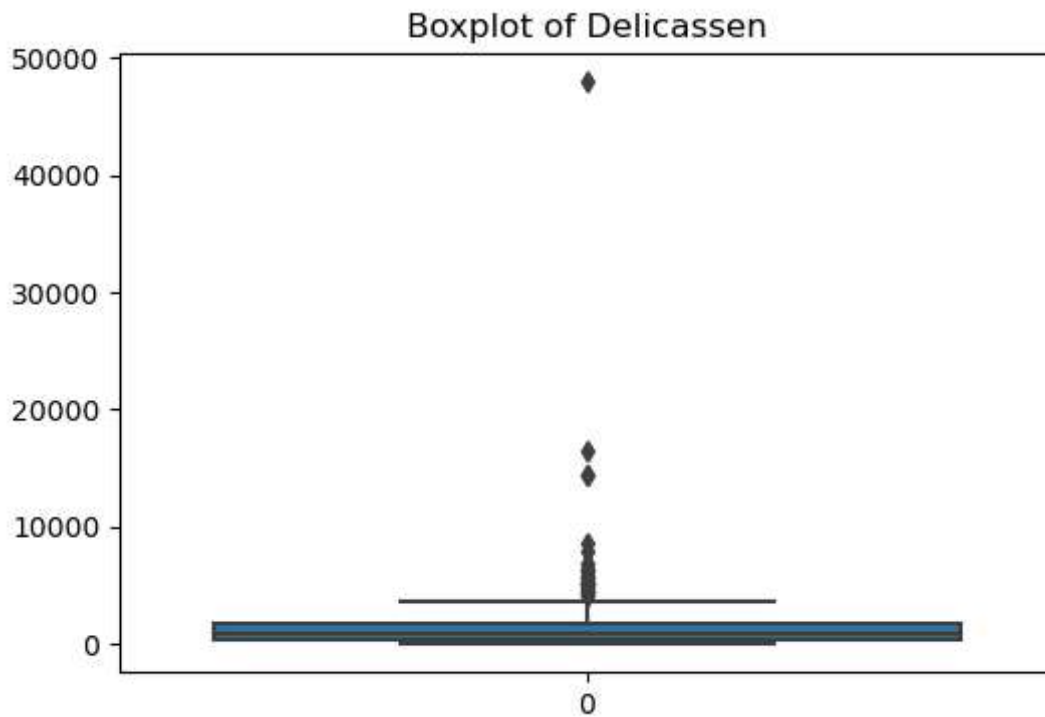
Boxplot of Milk



Boxplot of Grocery








Number of outliers in Channel: 0
Number of outliers in Region: 0
Number of outliers in Fresh: 20
Number of outliers in Milk: 28
Number of outliers in Grocery: 24
Number of outliers in Frozen: 43
Number of outliers in Detergents_Paper: 30
Number of outliers in Delicassen: 27

```
In [11]: def handle_outliers(dataframe, column):  
    Q1 = dataframe[column].quantile(0.25)  
    Q3 = dataframe[column].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_limit = Q1 - 1.5*IQR  
    upper_limit = Q3 + 1.5*IQR  
    dataframe[column] = dataframe[column].apply(lambda x: upper_limit if x > upper_limit else x)  
  
    for column in df.columns:  
        handle_outliers(df, column)
```


```
In [12]: import seaborn as sns
import matplotlib.pyplot as plt

for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.boxplot(df[column])
    plt.title(f'Boxplot of {column}')
    plt.show()


for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[column], bins=30, kde=True)
    plt.title(f'Distribution of {column}')
    plt.show()
```







Boxplot of Region





Boxplot of Fresh












Distribution of Region




Distribution of Fresh

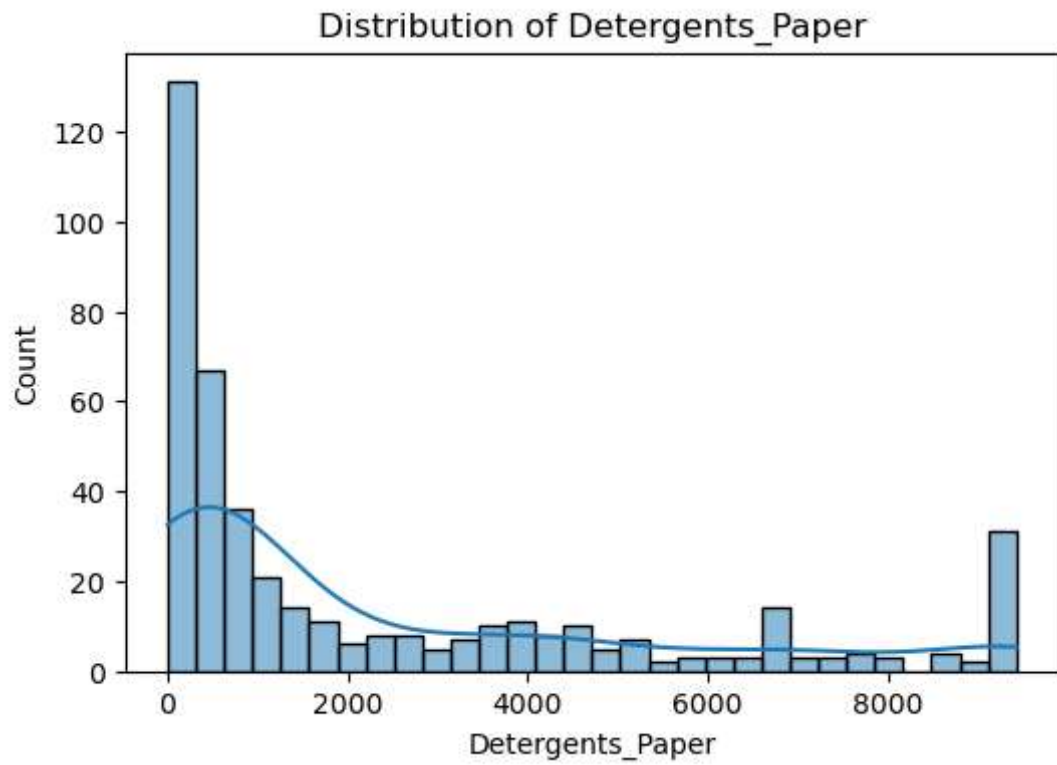




Distribution of Milk



Distribution of Grocery







```
In [13]: def detect_outliers(dataframe, column):
    Q1 = dataframe[column].quantile(0.25)
    Q3 = dataframe[column].quantile(0.75)
    IQR = Q3 - Q1
    outliers = dataframe[(dataframe[column] < Q1 - 1.5*IQR) | (dataframe[column] > Q3 + 1.5*IQR)]
    return outliers

for column in df.columns:
    outliers = detect_outliers(df, column)
    print(f'Number of outliers in {column}: {len(outliers)}')
```

```
Number of outliers in Channel: 0
Number of outliers in Region: 0
Number of outliers in Fresh: 0
Number of outliers in Milk: 0
Number of outliers in Grocery: 0
Number of outliers in Frozen: 0
Number of outliers in Detergents_Paper: 0
Number of outliers in Delicassen: 0
```



```
In [14]: print("Descriptive Statistics:")
print(df.describe())

print("Number of duplicate rows: ", df.duplicated().sum())

for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[column], bins=30, kde=True)
    plt.title(f'Distribution of {column}')
    plt.show()

plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Heatmap')
plt.show()
```


Descriptive Statistics:

	Channel	Region	Fresh	Milk	Grocery \
count	440.000000	440.000000	440.000000	440.000000	440.000000
mean	1.322727	2.543182	11357.568182	5048.592045	7236.37500
std	0.468052	0.774272	10211.542235	4386.377073	6596.53308
min	1.000000	1.000000	3.000000	55.000000	3.000000
25%	1.000000	2.000000	3127.750000	1533.000000	2153.00000
50%	1.000000	3.000000	8504.000000	3627.000000	4755.50000
75%	2.000000	3.000000	16933.750000	7190.250000	10655.75000
max	2.000000	3.000000	37642.750000	15676.125000	23409.87500


	Frozen	Detergents_Paper	Delicassen
count	440.000000	440.000000	440.000000
mean	2507.085795	2392.616477	1266.715341
std	2408.297738	2940.794090	1083.069792
min	25.000000	3.000000	3.000000
25%	742.250000	256.750000	408.250000
50%	1526.000000	816.500000	965.500000
75%	3554.250000	3922.000000	1820.250000
max	7772.250000	9419.875000	3938.250000

Number of duplicate rows: 0


Distribution of Channel




Distribution of Region




Distribution of Fresh




Distribution of Milk




Distribution of Grocery




Distribution of Frozen




Distribution of Detergents_Paper



Distribution of Delicassen





```
In [15]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)
```


```
In [16]: from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

wcss = []
max_clusters = 15
for i in range(1, max_clusters+1):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
    kmeans.fit(df)
    wcss.append(kmeans.inertia_)

plt.plot(range(1, max_clusters+1), wcss)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.grid(True)
plt.show()
```

[illegible]

reWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn()



```
In [17]: from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=4, init='k-means++', random_state=42)
kmeans.fit(df)
```

```
cluster_labels = kmeans.labels_
```

```
df['Cluster'] = cluster_labels
```

```
print(df.head())
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	\
0	2	3	12669.0	9656.0	7561.0	214.0	2674.0	
1	2	3	7057.0	9810.0	9568.0	1762.0	3293.0	
2	2	3	6353.0	8808.0	7684.0	2405.0	3516.0	
3	1	3	13265.0	1196.0	4221.0	6404.0	507.0	
4	2	3	22615.0	5410.0	7198.0	3915.0	1777.0	

	Delicassen	Cluster
0	1338.00	0
1	1776.00	1
2	3938.25	3
3	1788.00	0
4	3938.25	0

```
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
```

```
In [18]: df['Cluster'] = kmeans.labels_  
  
print("Cluster Sizes:\n", df['Cluster'].value_counts())  
  
for i in range(4):  
    print("\nCluster ", i)  
    print(df[df['Cluster'] == i].describe())
```

Cluster Sizes:

3 176

0 112

1 94

2 58

Name: Cluster, dtype: int64

Cluster 0

	Channel	Region	Fresh	Milk	Grocery \
count	112.000000	112.000000	112.000000	112.000000	112.000000
mean	1.214286	2.535714	16051.205357	3135.813616	4211.589286
std	0.412170	0.781873	3763.633078	2524.464860	3150.441587
min	1.000000	1.000000	10379.000000	134.000000	3.000000
25%	1.000000	2.000000	12419.750000	1283.500000	1970.500000
50%	1.000000	3.000000	16195.000000	2252.000000	3203.000000
75%	1.000000	3.000000	18830.250000	4537.000000	5700.250000
max	2.000000	3.000000	24929.000000	15676.125000	14982.000000

	Frozen	Detergents_Paper	Delicassen	Cluster
count	112.000000	112.000000	112.000000	112.0
mean	2988.859375	994.785714	1229.573661	0.0
std	2531.352938	1245.589613	963.527882	0.0
min	118.000000	3.000000	51.000000	0.0
25%	1018.750000	188.500000	514.250000	0.0
50%	2157.500000	456.500000	879.000000	0.0
75%	4276.000000	1404.000000	1804.500000	0.0
max	7772.250000	6707.000000	3938.250000	0.0

Cluster 1

	Channel	Region	Fresh	Milk	Grocery \
count	94.000000	94.000000	94.000000	94.000000	94.000000
mean	1.893617	2.489362	5331.893617	10454.450798	17196.140957
std	0.309980	0.799794	5111.448153	3937.245330	4905.345002
min	1.000000	1.000000	18.000000	1266.000000	8852.000000
25%	2.000000	2.000000	1409.500000	7576.000000	12563.250000
50%	2.000000	3.000000	4047.000000	10601.000000	16596.000000
75%	2.000000	3.000000	7870.500000	14316.500000	22288.500000
max	2.000000	3.000000	22925.000000	15676.125000	23409.875000

	Frozen	Detergents_Paper	Delicassen	Cluster
count	94.000000	94.000000	94.000000	94.0
mean	1496.428191	6936.898936	1547.364362	1.0
std	1538.882840	2383.035957	1176.131062	0.0
min	25.000000	241.000000	3.000000	1.0
25%	438.500000	5274.250000	680.000000	1.0
50%	973.000000	6931.500000	1366.500000	1.0
75%	1900.000000	9419.875000	2157.750000	1.0
max	7772.250000	9419.875000	3938.250000	1.0

Cluster 2

	Channel	Region	Fresh	Milk	Grocery \
count	58.000000	58.000000	58.000000	58.000000	58.000000
mean	1.172414	2.655172	32136.810345	5973.515086	7309.012931
std	0.381039	0.714554	5122.024937	4808.223223	5915.174661
min	1.000000	1.000000	22647.000000	286.000000	471.000000
25%	1.000000	3.000000	27207.500000	2393.000000	2726.250000
50%	1.000000	3.000000	31664.000000	4347.000000	5259.500000

75%	1.000000	3.000000	37642.750000	7829.500000	9344.000000
max	2.000000	3.000000	37642.750000	15676.125000	23409.875000

	Frozen	Detergents_Paper	Delicassen	Cluster
count	58.000000	58.000000	58.000000	58.0
mean	4170.017241	1417.426724	1967.702586	2.0
std	2841.060439	2055.702539	1267.507352	0.0
min	127.000000	10.000000	3.000000	2.0
25%	1370.750000	250.250000	1037.250000	2.0
50%	3662.000000	617.500000	1821.500000	2.0
75%	7772.250000	1428.000000	2910.250000	2.0
max	7772.250000	9419.875000	3938.250000	2.0

Cluster 3


	Channel	Region	Fresh	Milk	Grocery \
count	176.000000	176.000000	176.000000	176.000000	176.000000
mean	1.136364	2.539773	4741.261364	3073.790483	3817.880682
std	0.344153	0.777254	3072.006036	2492.137013	2790.348628
min	1.000000	1.000000	3.000000	55.000000	137.000000
25%	1.000000	2.000000	2116.000000	1109.000000	1739.250000
50%	1.000000	3.000000	4659.500000	2268.000000	2765.500000
75%	1.000000	3.000000	7369.250000	4394.250000	5494.500000
max	2.000000	3.000000	10290.000000	15676.125000	12400.000000


	Frozen	Detergents_Paper	Delicassen	Cluster
count	176.000000	176.000000	176.000000	176.0
mean	2192.274148	1176.454545	909.451705	3.0
std	2210.017535	1473.393792	872.339683	0.0
min	47.000000	5.000000	3.000000	3.0
25%	587.750000	216.500000	308.250000	3.0
50%	1310.000000	472.500000	674.500000	3.0
75%	2964.250000	1545.000000	1154.750000	3.0
max	7772.250000	7271.000000	3938.250000	3.0


```
In [19]: cluster_means = df.groupby('Cluster').mean()


cluster_means = cluster_means.transpose()


for feature in cluster_means.index:
    cluster_means.loc[feature].plot(kind='bar', figsize=(8,6))
    plt.title(feature)
    plt.ylabel('Mean Value')
    plt.xticks(ticks=range(4), labels=['Cluster 0', 'Cluster 1', 'Cluster 2',
    plt.show()
```




















```

In [20]: from sklearn.decomposition import PCA
import matplotlib.pyplot as plt


pca = PCA(n_components=2)
principalComponents = pca.fit_transform(df.drop('Cluster', axis=1))

PCA_components = pd.DataFrame(principalComponents, columns=['Principal Component 1', 'Principal Component 2'])

PCA_components['Cluster'] = df['Cluster']

plt.figure(figsize=(8,6))
plt.scatter(PCA_components['Principal Component 1'], PCA_components['Principal Component 2'])
plt.title('Clusters in PCA 2D Space')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster')
plt.show()

```





Conclusion:

1. The clusters obtained by implementing the Kmeans algorithm on the wholesale customers dataset can be used for a wide variety of business use cases such as grouping customers with similar purchasing behaviors together to make better marketing strategies or to recommend products to customers based on which cluster which they belong too or by exploring associations between products that are frequently purchased together within each cluster etc.
2. The customers present in each of the different clusters may have diverse needs and expectations and the impact of a specific delivery scheme can vary widely among them however for each customer segment by assessing how well a proposed delivery scheme aligns with their characteristics and preferences can be considered.